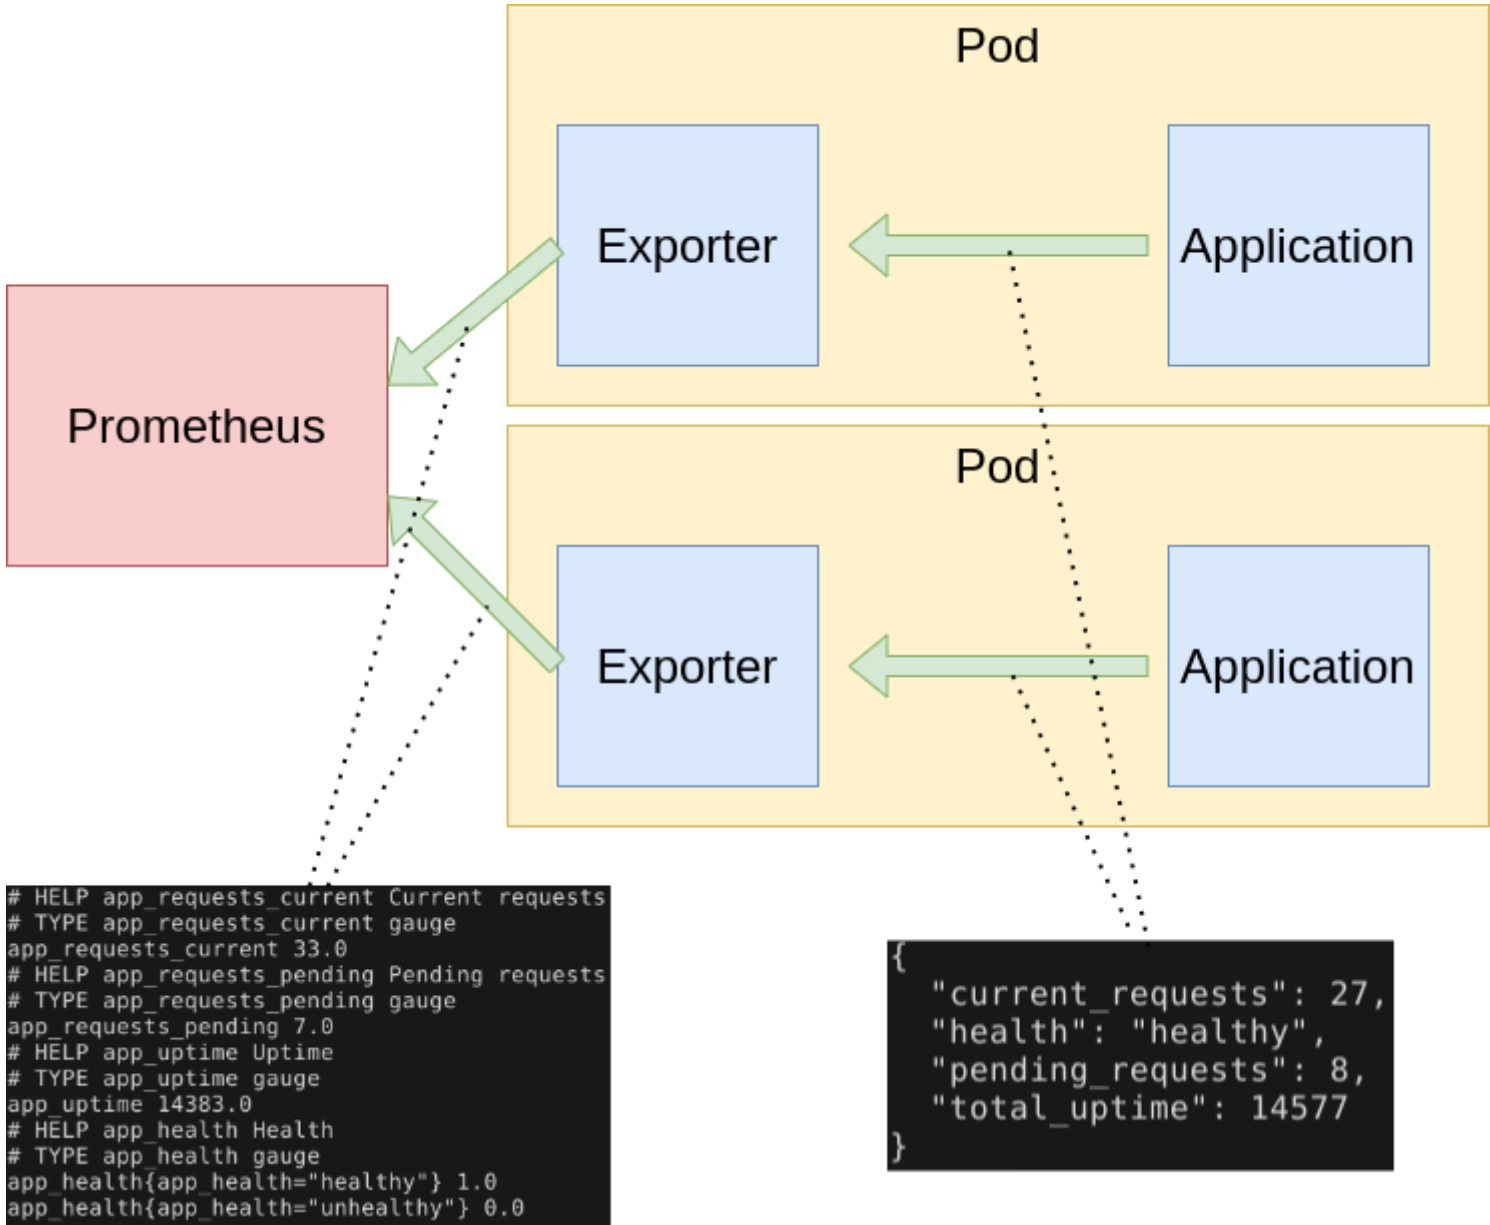


# Create a Quick and Easy Prometheus Exporter

Posted 1 year ago by Thomas Stringer

[Prometheus](#) is a clear leader in the cloud native world for metrics. Prometheus follows an HTTP pull model: It scrapes Prometheus metrics from endpoints routinely. Typically the abstraction layer between the application and Prometheus is an **exporter**, which takes application-formatted metrics and converts them to Prometheus metrics for consumption. Because Prometheus is an HTTP pull model, the exporter typically provides an endpoint where the Prometheus metrics can be scraped.

The relationship between Prometheus, the exporter, and the application in a Kubernetes environment can be visualized like this:



As you can see from above, the role of the exporter is to consume application-formatted metrics and transform them into Prometheus metrics. In the Kubernetes world, the exporter is a container that lives in the same pod as the application it is exporting.

## Do I need to write my own exporter?

You might need to write your own exporter if...

- You're using 3rd party software that doesn't have an [existing exporter](#) already
- You want to generate Prometheus metrics from software that you have written

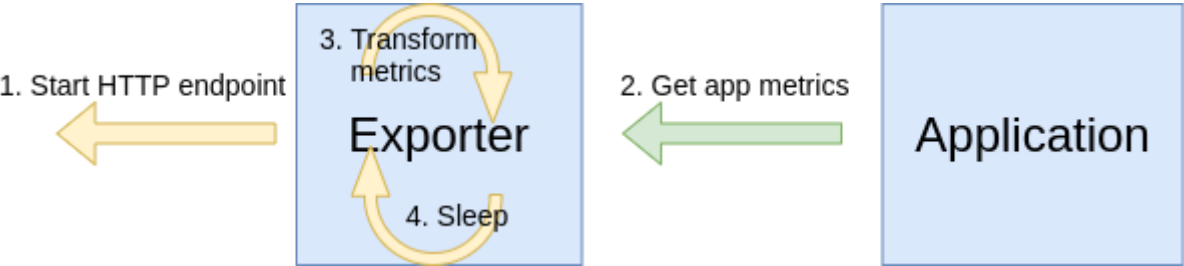
If you decide that you need to write your exporter, there are a handful of available clients to make things easier: [Python](#), [Go](#), [Java](#), and a [list of others](#).

## Using Python

As you can see above, there are a multiple languages and client libraries that you can use to create your exporter, but I've found that the Python approach is the quickest and easiest to get a working exporter.

is rapid development. So if you're in a DevOps shop that is tasked with writing exporters with a strict timeline, Python might be the best option for you.

The way an exporter works is shown below:



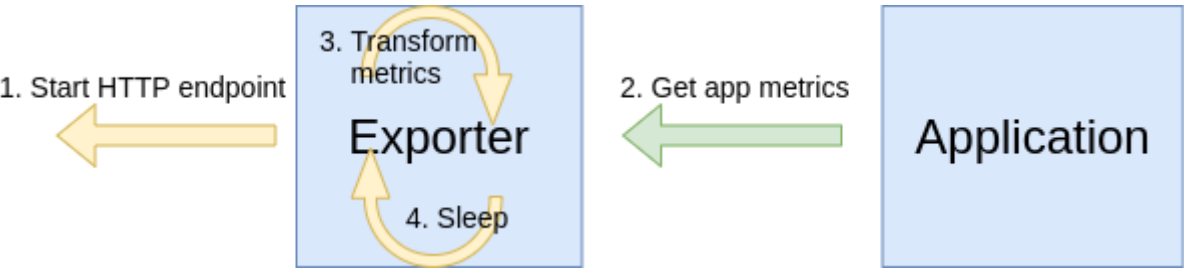
Steps 2, 3, and 4 happen in a loop.

Let's see a small but complete exporter implemented in Python (notes and explanation in comments and below the snippet):

**exporter.py**

```
2
3 import os
4 import time
5 from prometheus_client import start_http_server, Gauge, Enum
6 import requests
7
8 class AppMetrics:
9     """
10     Representation of Prometheus metrics and loop to fetch and transform
11     application metrics into Prometheus metrics.
12     """
13
14     def __init__(self, app_port=80, polling_interval_seconds=5):
15         self.app_port = app_port
16         self.polling_interval_seconds = polling_interval_seconds
17
18         # Prometheus metrics to collect
19         self.current_requests = Gauge("app_requests_current", "Current requests")
20         self.pending_requests = Gauge("app_requests_pending", "Pending requests")
21         self.total_uptime = Gauge("app_uptime", "Uptime")
22         self.health = Enum("app_health", "Health", states=["healthy", "unhealthy"])
23
24     def run_metrics_loop(self):
25         """Metrics fetching loop"""
26
27         while True:
28             self.fetch()
29             time.sleep(self.polling_interval_seconds)
30
31     def fetch(self):
32         """
33         Get metrics from application and refresh Prometheus metrics with
34         new values.
35         """
36
37         # Fetch raw status data from the application
38         resp = requests.get(url=f"http://localhost:{self.app_port}/status")
39         status_data = resp.json()
40
41         # Update Prometheus metrics with application metrics
42         self.current_requests.set(status_data["current_requests"])
43         self.pending_requests.set(status_data["pending_requests"])
44         self.total_uptime.set(status_data["total_uptime"])
45         self.health.state(status_data["health"])
46
47 def main():
48     """Main entry point"""
49
50     polling_interval_seconds = int(os.getenv("POLLING_INTERVAL_SECONDS", "5"))
51     app_port = int(os.getenv("APP_PORT", "80"))
52     exporter_port = int(os.getenv("EXPORTER_PORT", "9877"))
53
54     app_metrics = AppMetrics(
55         app_port=app_port,
56         polling_interval_seconds=polling_interval_seconds
57     )
58     start_http_server(exporter_port)
59     app_metrics.run_metrics_loop()
60
61 if __name__ == "__main__":
62     main()
```

class wasn't necessary, but I think it's a bit cleaner and more future proof to start this way and keep the logic and data contained in an object. Effectively the code just loops and makes an HTTP request to the application metrics (available on `localhost:APP_PORT/status` ) and transforms that to the Prometheus metrics.



- Step 1 is the call to `start_http_server` in `main` (line 58)
- Step 2 is `AppMetrics.fetch` (line 31), which is invoked from the loop implemented in `AppMetrics.run_metrics_loop` (line 24)
- Step 3 is the multiple calls `set` and `state` calls on the Prometheus metrics in `AppMetrics.fetch` (lines 42 - 45)
- Step 4, the sleep, is implemented in `AppMetrics.run_metrics_loop` so that we have some delay in metrics scraping

In Kubernetes, this could be deployed alongside the application container like this:

deployment.yaml

```
1 kind: Deployment
2 apiVersion: apps/v1
3 metadata:
4   name: webapp
5 spec:
6   replicas: 8
7   selector:
8     matchLabels:
9       app: webapp
10  template:
11    metadata:
12      labels:
13        app: webapp
14    spec:
15      containers:
16        - name: webapp
17          image: mycontainerregistry/webapp:latest
18          imagePullPolicy: Always
19          ports:
20            - containerPort: 5000
21              name: http
22        - name: exporter
23          image: mycontainerregistry/webappexporter:latest
24          imagePullPolicy: Always
25          env:
26            - name: POLLING_INTERVAL_SECONDS
27              value: "5"
28            - name: APP_PORT
29              value: "5000"
30            - name: EXPORTER_PORT
31              value: "9877"
32          ports:
33            - containerPort: 9877
34              name: http
```

Now with this exporter, Prometheus can scrape metrics from this pod on port 9877, all thanks to the exporter shim that was put between Prometheus and the application!

More on exporters

[with the different types of Prometheus metrics](#) available.


## Summary

Prometheus is a really great and powerful tool. At first glance, it could seem like a daunting task to create an exporter, but hopefully this blog post has shown that in just a few lines of Python code you can instrument an effective exporter and start pulling your application metrics right away into Prometheus!

 [Blog](#)

 [prometheus](#) [kubernetes](#) [devops](#) [python](#)

This post is licensed under [CC BY 4.0](#) by the author.

Share:    

OLDER

[Why You Should Typically Use Named Arguments in Python](#)

NEWER

[Debug a Python Application Running in Kubernetes](#)

## Further Reading

[1 year ago](#)

[Collect Custom Metrics in AKS](#)

[Part of running a production-quality Kubernetes cluster is being able to monit...](#)

[1 year ago](#)

[Beware of kubectl's -t \(--tty\) Option](#)

[I see this all over the place in documentation, blog posts, scripts,...](#)

[1 year ago](#)

[Run Kubernetes in Azure the Cheap Way](#)

[Update: Since writing this blog post, I have found that you can save another few...](#)